

# Model learning with structured latent representations

---

Sharad Vikram

June 21, 2017

UCSD

# Background

---

## Background

- Probabilistic graphical models

- Model learning

- Variational inference

- Structured variational autoencoder

- Model learning

## Conclusion

- Applications

- Current work

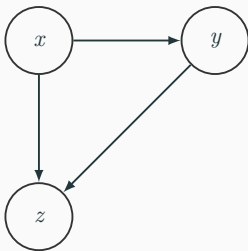
# Probabilistic graphical models

A probabilistic graphical model (PGM) is a graphical representation of a joint distribution over several random variables.

# Probabilistic graphical models

A probabilistic graphical model (PGM) is a graphical representation of a joint distribution over several random variables.

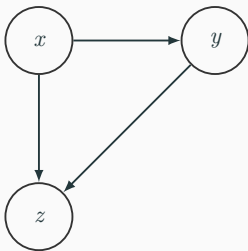
For example, a possible graph for variables  $x$ ,  $y$  and  $z$  is



# Probabilistic graphical models

A probabilistic graphical model (PGM) is a graphical representation of a joint distribution over several random variables.

For example, a possible graph for variables  $x$ ,  $y$  and  $z$  is



The joint distribution factorizes as

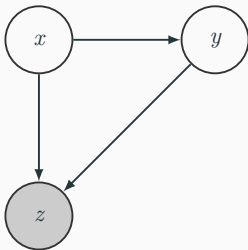
$$p(x, y, z) = p(x)p(y|x)p(z|x, y)$$

# Bayesian inference

Bayesian inference allows us to compute probability distributions after observations have been made.

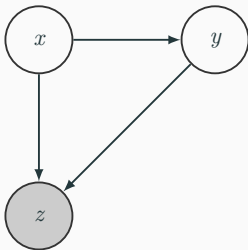
# Bayesian inference

Bayesian inference allows us to compute probability distributions after observations have been made.



# Bayesian inference

Bayesian inference allows us to compute probability distributions after observations have been made.

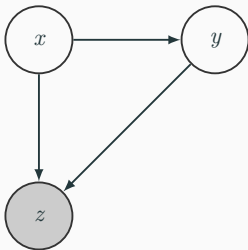


We are interested in the posterior distribution  $p(x, y|z)$



# Bayesian inference

Bayesian inference allows us to compute probability distributions after observations have been made.




We are interested in the posterior distribution  $p(x, y|z)$

This can be computed via Bayes rule:

$$p(x, y|z) = \frac{p(x, y, z)}{p(z)} = \frac{p(x)p(y|x)p(z|x, y)}{\int p(x)p(y|x)p(z|x, y) dx dy}$$


# Example PGM

Latent variable model: Gaussian mixture model



# Example PGM


Latent variable model: Gaussian mixture model



Local variables:  $\{z_i\}_{i=1}^N$

# Example PGM

Latent variable model: Gaussian mixture model



Local variables:  $\{z_i\}_{i=1}^N$

Global variables:  $\mu, \pi$

# Conjugacy

## Conjugate

Two random variables  $x$  and  $y$  whose distribution is

$$p(x, y) = p(x)p(y|x)$$

are said to be *conjugate* if the posterior  $p(x|y)$  is in the same family of distributions as  $p(x)$ .

# Conjugacy

## Conjugate

Two random variables  $x$  and  $y$  whose distribution is

$$p(x, y) = p(x)p(y|x)$$

are said to be *conjugate* if the posterior  $p(x|y)$  is in the same family of distributions as  $p(x)$ .

Examples of conjugate distributions:

- Normal/normal

# Conjugacy

## Conjugate

Two random variables  $x$  and  $y$  whose distribution is

$$p(x, y) = p(x)p(y|x)$$

are said to be *conjugate* if the posterior  $p(x|y)$  is in the same family of distributions as  $p(x)$ .

Examples of conjugate distributions:

- Normal/normal
- Normal-inverse-wishart(NIW)/normal

# Conjugacy

## Conjugate

Two random variables  $x$  and  $y$  whose distribution is

$$p(x, y) = p(x)p(y|x)$$

are said to be *conjugate* if the posterior  $p(x|y)$  is in the same family of distributions as  $p(x)$ .

Examples of conjugate distributions:

- Normal/normal
- Normal-inverse-wishart(NIW)/normal
- Dirichlet/multinomial



# Algorithms for inference

Depending on the structure of the graphical model and conjugacy of the variables, computing the posterior distribution can be easy, or very hard!

# Algorithms for inference

Depending on the structure of the graphical model and conjugacy of the variables, computing the posterior distribution can be easy, or very hard!

Possible ways:

- Compute posterior analytically

# Algorithms for inference

Depending on the structure of the graphical model and conjugacy of the variables, computing the posterior distribution can be easy, or very hard!

Possible ways:

- Compute posterior analytically
- Sampling (MCMC, Gibbs sampling)

# Algorithms for inference

Depending on the structure of the graphical model and conjugacy of the variables, computing the posterior distribution can be easy, or very hard!

Possible ways:

- Compute posterior analytically
- Sampling (MCMC, Gibbs sampling)
- Expectation-maximization

# Algorithms for inference

Depending on the structure of the graphical model and conjugacy of the variables, computing the posterior distribution can be easy, or very hard!

Possible ways:

- Compute posterior analytically
- Sampling (MCMC, Gibbs sampling)
- Expectation-maximization
- Variational inference

# Algorithms for inference

Depending on the structure of the graphical model and conjugacy of the variables, computing the posterior distribution can be easy, or very hard!


Possible ways:

- Compute posterior analytically
- Sampling (MCMC, Gibbs sampling)
- Expectation-maximization
- Variational inference

PGMs offer interpretability and quantifiable uncertainty, but often don't scale well with data and can be underexpressive.


# Model learning

Consider an agent in a system, and a set of its possible actions.



# Model learning

Consider an agent in a system, and a set of its possible actions.



Model learning is the problem of estimating the dynamics of the system when we don't know it beforehand.



Formally, consider an agent in a system with state space  $\mathcal{S}$  with action space  $\mathcal{A}$  with underlying dynamics function  $p(s_{t+1}|s_t, a_t)$ .

We are interested in learning an approximate dynamics function

$$\hat{p}(s_{t+1}|s_t, a_t)$$

from a dataset of trajectories

$$\tau = \{(s_0^{(i)}, a_0^{(i)}, s_1^{(i)}, a_1^{(i)}, \dots, s_T^{(i)})\}_{i=1}^N$$

# Bayesian linear dynamical system

A simple assumption: **Bayesian linear dynamical system** (LDS)

$$\mu_\rho, \Sigma_\rho \sim \mathcal{NIW}(\Psi, \nu, \mu_0, \kappa), \quad \mathbf{F}, \Sigma \sim \mathcal{MNIW}(\Psi, \nu, \mathbf{M}_0, \mathbf{V}),$$


$$\mathbf{s}_0 \mid \mu_\rho, \Sigma_\rho \sim \mathcal{N}(\mu_\rho, \Sigma_\rho), \quad \mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t \sim \mathcal{N} \left( \mathbf{F} \begin{bmatrix} \mathbf{s}_t \\ \mathbf{a}_t \end{bmatrix}, \Sigma \right) \text{ for } t \in [0, \dots, T]$$

# Bayesian linear dynamical system

A simple assumption: **Bayesian linear dynamical system** (LDS)

$$\mu_\rho, \Sigma_\rho \sim \mathcal{NIW}(\Psi, \nu, \mu_0, \kappa), \quad \mathbf{F}, \Sigma \sim \mathcal{MNIW}(\Psi, \nu, \mathbf{M}_0, \mathbf{V}),$$

$$\mathbf{s}_0 \mid \mu_\rho, \Sigma_\rho \sim \mathcal{N}(\mu_\rho, \Sigma_\rho), \quad \mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t \sim \mathcal{N}\left(\mathbf{F} \begin{bmatrix} \mathbf{s}_t \\ \mathbf{a}_t \end{bmatrix}, \Sigma\right) \text{ for } t \in [0, \dots, T]$$




# Bayesian linear dynamical system

A simple assumption: **Bayesian linear dynamical system** (LDS)

$$\mu_\rho, \Sigma_\rho \sim \mathcal{NIW}(\Psi, \nu, \mu_0, \kappa), \quad \mathbf{F}, \Sigma \sim \mathcal{MNIW}(\Psi, \nu, \mathbf{M}_0, \mathbf{V}),$$

$$\mathbf{s}_0 \mid \mu_\rho, \Sigma_\rho \sim \mathcal{N}(\mu_\rho, \Sigma_\rho), \quad \mathbf{s}_{t+1} \mid \mathbf{s}_t, \mathbf{a}_t \sim \mathcal{N}\left(\mathbf{F} \begin{bmatrix} \mathbf{s}_t \\ \mathbf{a}_t \end{bmatrix}, \Sigma\right) \text{ for } t \in [0, \dots, T]$$



We are interested in the posterior distribution

$p(\mu_\rho, \Sigma_\rho, \mathbf{F}, \Sigma \mid s_0, a_0, \dots, s_T)$  which can be computed analytically.

# Variational inference at a high level

# Variational inference at a high level

Consider a latent variable model with global variables  $\theta$ , local variables  $z$  and observations  $x$ . Our desired posterior is  $p(\theta, z|x)$

# Variational inference at a high level

Consider a latent variable model with global variables  $\theta$ , local variables  $z$  and observations  $x$ . Our desired posterior is  $p(\theta, z|x)$

**Strategy:** convert inference into optimization

# Variational inference at a high level

Consider a latent variable model with global variables  $\theta$ , local variables  $z$  and observations  $x$ . Our desired posterior is  $p(\theta, z|x)$

**Strategy:** convert inference into optimization

- Instantiate *variational distribution*  $q_{\phi}(\theta, z)$  where  $\phi$  are free parameters



# Variational inference at a high level

Consider a latent variable model with global variables  $\theta$ , local variables  $z$  and observations  $x$ . Our desired posterior is  $p(\theta, z|x)$

**Strategy:** convert inference into optimization

- Instantiate *variational distribution*  $q_\phi(\theta, z)$  where  $\phi$  are free parameters
- Define loss  $\text{KL}(q_\phi(\theta, z) \| p(\theta, z|x))$

# Variational inference at a high level

Consider a latent variable model with global variables  $\theta$ , local variables  $z$  and observations  $x$ . Our desired posterior is  $p(\theta, z|x)$

**Strategy:** convert inference into optimization

- Instantiate *variational distribution*  $q_\phi(\theta, z)$  where  $\phi$  are free parameters
- Define loss  $\text{KL}(q_\phi(\theta, z) \| p(\theta, z|x))$
- Minimize loss  $\phi^* = \operatorname{argmin}_\phi \text{KL}(q_\phi(\theta, z) \| p(\theta, z|x))$

# Variational inference at a high level

Consider a latent variable model with global variables  $\theta$ , local variables  $z$  and observations  $x$ . Our desired posterior is  $p(\theta, z|x)$

**Strategy:** convert inference into optimization

- Instantiate *variational distribution*  $q_\phi(\theta, z)$  where  $\phi$  are free parameters
- Define loss  $\text{KL}(q_\phi(\theta, z)||p(\theta, z|x))$
- Minimize loss  $\phi^* = \operatorname{argmin}_\phi \text{KL}(q_\phi(\theta, z)||p(\theta, z|x))$

If  $q(\theta, z)$  is sufficiently expressive, it can approximate  $p(\theta, z|x)$  quite well.

# KL-divergence

Kullback-Leibler (KL) divergence is a measure of how far one probability distribution is from another.

# KL-divergence

Kullback-Leibler (KL) divergence is a measure of how far one probability distribution is from another.

For distributions  $q(x)$  and  $p(x)$ ,

$$\text{KL}(q(x) \| p(x)) = \int q(x) \log \frac{q(x)}{p(x)} dx$$

# KL-divergence

Kullback-Leibler (KL) divergence is a measure of how far one probability distribution is from another.

For distributions  $q(x)$  and  $p(x)$ ,

$$\text{KL}(q(x)||p(x)) = \int q(x) \log \frac{q(x)}{p(x)} dx$$

**Properties:**

- $\text{KL}(q(x)||p(x)) = 0$  if  $q(x) = p(x)$ .
- Asymmetric

## Evidence lower bound

In general, we cannot even compute  $KL(q(\theta, z) \| p(\theta, z|x))$  because we don't know the posterior  $p(\theta, z|x)$ .

## Evidence lower bound

In general, we cannot even compute  $KL(q(\theta, z)||p(\theta, z|x))$  because we don't know the posterior  $p(\theta, z|x)$ .

We can rewrite the KL divergence as

$$\begin{aligned} KL(q(\theta, z)||p(\theta, z|x)) &= \int q(\theta, z) \log \frac{q(\theta, z)}{p(\theta, z|x)} d\theta, z \\ &= \log p(x) - \mathbb{E}_{q(\theta, z)} \left[ \log \frac{p(x, \theta, z)}{q(\theta, z)} \right] \end{aligned}$$



## Evidence lower bound

In general, we cannot even compute  $KL(q(\theta, z) \| p(\theta, z|x))$  because we don't know the posterior  $p(\theta, z|x)$ .

We can rewrite the KL divergence as

$$\begin{aligned} KL(q(\theta, z) \| p(\theta, z|x)) &= \int q(\theta, z) \log \frac{q(\theta, z)}{p(\theta, z|x)} d\theta, z \\ &= \log p(x) - \mathbb{E}_{q(\theta, z)} \left[ \log \frac{p(x, \theta, z)}{q(\theta, z)} \right] \end{aligned}$$

and maximize the *evidence lower bound* (ELBO)

$$\mathcal{L}[q(\theta, z)] = \mathbb{E}_{q(\theta, z)} \left[ \log \frac{p(x, \theta, z)}{q(\theta, z)} \right]$$

# Variational inference for PGMs

For a general graphical model with variable set  $\mathbf{X} = \{x_1, x_2, \dots\}$  we have joint distribution

$$p(\mathbf{X}) = \prod_i p(x_i | \text{pa}_i)$$

where  $\text{pa}_i$  are the parents of node  $x_i$  in the graph.

# Variational inference for PGMs

For a general graphical model with variable set  $\mathbf{X} = \{x_1, x_2, \dots\}$  we have joint distribution

$$p(\mathbf{X}) = \prod_i p(x_i | \text{pa}_i)$$

where  $\text{pa}_i$  are the parents of node  $x_i$  in the graph.

Let the set  $\mathbf{H}$  be all unobserved variables and  $\mathbf{V}$  be the observed.

# Variational inference for PGMs

For a general graphical model with variable set  $\mathbf{X} = \{x_1, x_2, \dots\}$  we have joint distribution

$$p(\mathbf{X}) = \prod_i p(x_i | \text{pa}_i)$$

where  $\text{pa}_i$  are the parents of node  $x_i$  in the graph.

Let the set  $\mathbf{H}$  be all unobserved variables and  $\mathbf{V}$  be the observed.

We are interested in the posterior  $p(\mathbf{H} | \mathbf{V})$  and use variational distribution

$$q(\mathbf{H}) = \prod_i q(\mathbf{H}_i)$$

# Variational inference for PGMs

For a general graphical model with variable set  $\mathbf{X} = \{x_1, x_2, \dots\}$  we have joint distribution

$$p(\mathbf{X}) = \prod_i p(x_i | \text{pa}_i)$$

where  $\text{pa}_i$  are the parents of node  $x_i$  in the graph.

Let the set  $\mathbf{H}$  be all unobserved variables and  $\mathbf{V}$  be the observed.

We are interested in the posterior  $p(\mathbf{H} | \mathbf{V})$  and use variational distribution

$$q(\mathbf{H}) = \prod_i q(\mathbf{H}_i)$$

This is called the *mean-field* assumption.

## Optimizing a single factor

We can isolate a single factor for each hidden node in the graph  $\mathbf{H}_j$ .  
This makes optimizing a single variational factor easy!

## Optimizing a single factor

We can isolate a single factor for each hidden node in the graph  $\mathbf{H}_j$ .  
This makes optimizing a single variational factor easy!

$$\mathcal{L}[q(\mathbf{H})] = \text{KL}(q(\mathbf{H}_j) \| \tilde{q}(\mathbf{H}_j, \mathbf{V})) + \text{const.}$$

# Optimizing a single factor

We can isolate a single factor for each hidden node in the graph  $\mathbf{H}_j$ . This makes optimizing a single variational factor easy!

$$\mathcal{L}[q(\mathbf{H})] = \text{KL}(q(\mathbf{H}_j) \parallel \tilde{q}(\mathbf{H}_j, \mathbf{V})) + \text{const.}$$

For a single factor  $q(\mathbf{H}_j)$ , this equation is minimized when  $q(\mathbf{H}_j) = \tilde{q}(\mathbf{H}_j, \mathbf{V})$ .



# Optimizing a single factor

We can isolate a single factor for each hidden node in the graph  $\mathbf{H}_j$ . This makes optimizing a single variational factor easy!

$$\mathcal{L}[q(\mathbf{H})] = \text{KL}(q(\mathbf{H}_j) \| \tilde{q}(\mathbf{H}_j, \mathbf{V})) + \text{const.}$$

For a single factor  $q(\mathbf{H}_j)$ , this equation is minimized when  $q(\mathbf{H}_j) = \tilde{q}(\mathbf{H}_j, \mathbf{V})$ .

Furthermore,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

is a function of only factors other than  $q(\mathbf{H}_j)$  and observed data.

# Optimizing a single factor

We can isolate a single factor for each hidden node in the graph  $\mathbf{H}_j$ . This makes optimizing a single variational factor easy!

$$\mathcal{L}[q(\mathbf{H})] = \text{KL}(q(\mathbf{H}_j) \parallel \tilde{q}(\mathbf{H}_j, \mathbf{V})) + \text{const.}$$

For a single factor  $q(\mathbf{H}_j)$ , this equation is minimized when  $q(\mathbf{H}_j) = \tilde{q}(\mathbf{H}_j, \mathbf{V})$ .

Furthermore,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

is a function of only factors other than  $q(\mathbf{H}_j)$  and observed data.

## Mean-field variational inference

Until converged, for each factor  $q(\mathbf{H}_j)$ , hold factors  $q(\mathbf{H}_{i \neq j})$  constant and set  $q(\mathbf{H}_j) = \tilde{q}(\mathbf{H}_j, \mathbf{V})$ .

# Conjugate-exponential graphical models

If our PGM is *conjugate-exponential*, where every node belongs in the exponential family of distributions, and is conjugate w.r.t. its parents,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

can be calculated in closed form!

# Conjugate-exponential graphical models

If our PGM is *conjugate-exponential*, where every node belongs in the exponential family of distributions, and is conjugate w.r.t. its parents,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

can be calculated in closed form!

An exponential family distribution is one that can be written in the following form:

$$p(x|\theta) = h(x) \exp \{ \langle \eta_x(\theta), t_x(x) \rangle - \log Z(\eta(\theta)) \}$$

with

# Conjugate-exponential graphical models

If our PGM is *conjugate-exponential*, where every node belongs in the exponential family of distributions, and is conjugate w.r.t. its parents,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

can be calculated in closed form!

An exponential family distribution is one that can be written in the following form:

$$p(x|\theta) = h(x) \exp \{ \langle \eta_x(\theta), t_x(x) \rangle - \log Z(\eta(\theta)) \}$$

with

- $h(x)$ : base measure

# Conjugate-exponential graphical models

If our PGM is *conjugate-exponential*, where every node belongs in the exponential family of distributions, and is conjugate w.r.t. its parents,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

can be calculated in closed form!

An exponential family distribution is one that can be written in the following form:

$$p(x|\theta) = h(x) \exp \{ \langle \eta_x(\theta), t_x(x) \rangle - \log Z(\eta(\theta)) \}$$

with

- $h(x)$ : base measure
- $\eta_x(\theta)$ : natural parameter

# Conjugate-exponential graphical models

If our PGM is *conjugate-exponential*, where every node belongs in the exponential family of distributions, and is conjugate w.r.t. its parents,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

can be calculated in closed form!

An exponential family distribution is one that can be written in the following form:

$$p(x|\theta) = h(x) \exp \{ \langle \eta_x(\theta), t_x(x) \rangle - \log Z(\eta(\theta)) \}$$

with

- $h(x)$ : base measure
- $\eta_x(\theta)$ : natural parameter
- $t_x(x)$ : sufficient statistic

# Conjugate-exponential graphical models

If our PGM is *conjugate-exponential*, where every node belongs in the exponential family of distributions, and is conjugate w.r.t. its parents,

$$\log \tilde{q}(\mathbf{H}_j, \mathbf{V}) = \mathbb{E}_{i \neq j} [\log p(\mathbf{H}, \mathbf{V})] + \text{const.}$$

can be calculated in closed form!

An exponential family distribution is one that can be written in the following form:

$$p(x|\theta) = h(x) \exp \{ \langle \eta_x(\theta), t_x(x) \rangle - \log Z(\eta(\theta)) \}$$


with

- $h(x)$ : base measure
- $\eta_x(\theta)$ : natural parameter
- $t_x(x)$ : sufficient statistic
- $\log Z(\eta_x(\theta))$ : log-partition function




# Variational message passing

We now return to graphs! If we assume a mean-field variational distribution and our PGM is conjugate-exponential, we get a very elegant graph algorithm, called *variational message passing* (VMP) [1].




# Variational message passing

We now return to graphs! If we assume a mean-field variational distribution and our PGM is conjugate-exponential, we get a very elegant graph algorithm, called *variational message passing* (VMP) [1].




# Variational message passing

We now return to graphs! If we assume a mean-field variational distribution and our PGM is conjugate-exponential, we get a very elegant graph algorithm, called *variational message passing* (VMP) [1].



# Variational message passing

We now return to graphs! If we assume a mean-field variational distribution and our PGM is conjugate-exponential, we get a very elegant graph algorithm, called *variational message passing* (VMP) [1].



Remember that  $\eta_j$  is the natural parameter of the distribution  $q(\mathbf{H}_j)$ .

Remember that  $\eta_j$  is the natural parameter of the distribution  $q(\mathbf{H}_j)$ .  
The message from a parent  $Y$  to a child  $X$  is

$$m_{Y \rightarrow X} = f(\eta_Y)$$

Remember that  $\eta_j$  is the natural parameter of the distribution  $q(\mathbf{H}_j)$ .  
The message from a parent  $Y$  to a child  $X$  is

$$m_{Y \rightarrow X} = f(\eta_Y)$$

The message from a child  $X$  to a parent  $Y$  is

$$m_{X \rightarrow Y} = g(\eta_X, \{m_{i \rightarrow X}\}_{i \in \text{cp}_Y})$$

Remember that  $\eta_j$  is the natural parameter of the distribution  $q(\mathbf{H}_j)$ .  
The message from a parent  $Y$  to a child  $X$  is

$$m_{Y \rightarrow X} = f(\eta_Y)$$

The message from a child  $X$  to a parent  $Y$  is

$$m_{X \rightarrow Y} = g(\eta_X, \{m_{i \rightarrow X}\}_{i \in \text{cp}_Y})$$

In conjugate-exponential PGMs, messages can be computed in closed-form.



**Setup:** Conjugate-exponential graphical model

# Summary of VMP

**Setup:** Conjugate-exponential graphical model

**Problem:** Compute posterior  $p(\mathbf{H}|\mathbf{V})$ , approximated with  
 $q(\mathbf{H}) = \prod_j q(\mathbf{H}_j)$

# Summary of VMP

**Setup:** Conjugate-exponential graphical model

**Problem:** Compute posterior  $p(\mathbf{H}|\mathbf{V})$ , approximated with  $q(\mathbf{H}) = \prod_j q(\mathbf{H}_j)$

**Solution:** Until converged, for each hidden node  $\mathbf{H}_j$ :

# Summary of VMP

**Setup:** Conjugate-exponential graphical model

**Problem:** Compute posterior  $p(\mathbf{H}|\mathbf{V})$ , approximated with  $q(\mathbf{H}) = \prod_j q(\mathbf{H}_j)$

**Solution:** Until converged, for each hidden node  $\mathbf{H}_j$ :

1. Collect messages from children and parents

# Summary of VMP

**Setup:** Conjugate-exponential graphical model

**Problem:** Compute posterior  $p(\mathbf{H}|\mathbf{V})$ , approximated with  $q(\mathbf{H}) = \prod_j q(\mathbf{H}_j)$

**Solution:** Until converged, for each hidden node  $\mathbf{H}_j$ :

1. Collect messages from children and parents
2. Compute updated distribution parameters from messages

# Summary of VMP

**Setup:** Conjugate-exponential graphical model

**Problem:** Compute posterior  $p(\mathbf{H}|\mathbf{V})$ , approximated with  $q(\mathbf{H}) = \prod_j q(\mathbf{H}_j)$

**Solution:** Until converged, for each hidden node  $\mathbf{H}_j$ :

1. Collect messages from children and parents
2. Compute updated distribution parameters from messages

**Benefits:** efficient, simple, can incorporate mini-batches (stochastic variational inference)

# Summary of VMP

**Setup:** Conjugate-exponential graphical model

**Problem:** Compute posterior  $p(\mathbf{H}|\mathbf{V})$ , approximated with  $q(\mathbf{H}) = \prod_j q(\mathbf{H}_j)$

**Solution:** Until converged, for each hidden node  $\mathbf{H}_j$ :

1. Collect messages from children and parents
2. Compute updated distribution parameters from messages

**Benefits:** efficient, simple, can incorporate mini-batches (stochastic variational inference)

**Drawbacks:** can be underexpressive (conjugate-exponential requirement)

# Structured variational autoencoder

---

## Background

- Probabilistic graphical models

- Model learning

- Variational inference

## Structured variational autoencoder

- Model learning

## Conclusion

- Applications

- Current work




# Model learning

Recall the model learning problem.


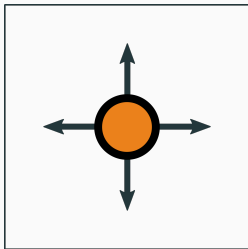
# Model learning

Recall the model learning problem.



# Model learning

Recall the model learning problem.




## Adding expressivity

One way of making the Bayesian LDS more expressive is to add a observation model.


# Adding expressivity

One way of making the Bayesian LDS more expressive is to add a observation model.



# Adding expressivity

One way of making the Bayesian LDS more expressive is to add a observation model.



What if this observation model was a neural network?

# Structured variational autoencoder

We augment the Bayesian LDS with a neural network observation model.

# Structured variational autoencoder

We augment the Bayesian LDS with a neural network observation model.

$$\begin{aligned}\mu_\rho, \Sigma_\rho &\sim \mathcal{NIW}(\Psi, \nu, \mu_0, \kappa), \quad \mathbf{F}, \Sigma \sim \mathcal{MNIW}(\Psi, \nu, \mathbf{M}_0, \mathbf{V}), \\ \mathbf{x}_0 \mid \mu_\rho, \Sigma_\rho &\sim \mathcal{N}(\mu_\rho, \Sigma_\rho), \quad \mathbf{x}_{t+1} \mid \mathbf{x}_t, \mathbf{a}_t \sim \mathcal{N}\left(\mathbf{F} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{a}_t \end{bmatrix}, \Sigma\right) \text{ for } t \in [0, \dots, T] \\ \mathbf{s}_t \mid \mathbf{x}_t &\sim \mathcal{N}(\mu_\gamma(\mathbf{x}_t), \Sigma_\gamma(\mathbf{x}_t)) \text{ for } t \in [0, \dots, T]\end{aligned}$$

where  $\mu_\gamma(\mathbf{x}_t)$  and  $\Sigma_\gamma(\mathbf{x}_t)$  are both neural networks parametrized by  $\gamma$ . This model is called a structured variational autoencoder (SVAE) [2].



# Structured variational autoencoder

We augment the Bayesian LDS with a neural network observation model.

$$\begin{aligned}\mu_\rho, \Sigma_\rho &\sim \mathcal{NITW}(\Psi, \nu, \mu_0, \kappa), \quad \mathbf{F}, \Sigma \sim \mathcal{MNITW}(\Psi, \nu, \mathbf{M}_0, \mathbf{V}), \\ \mathbf{x}_0 \mid \mu_\rho, \Sigma_\rho &\sim \mathcal{N}(\mu_\rho, \Sigma_\rho), \quad \mathbf{x}_{t+1} \mid \mathbf{x}_t, \mathbf{a}_t \sim \mathcal{N}\left(\mathbf{F} \begin{bmatrix} \mathbf{x}_t \\ \mathbf{a}_t \end{bmatrix}, \Sigma\right) \text{ for } t \in [0, \dots, T] \\ \mathbf{s}_t \mid \mathbf{x}_t &\sim \mathcal{N}(\mu_\gamma(\mathbf{x}_t), \Sigma_\gamma(\mathbf{x}_t)) \text{ for } t \in [0, \dots, T]\end{aligned}$$

where  $\mu_\gamma(\mathbf{x}_t)$  and  $\Sigma_\gamma(\mathbf{x}_t)$  are both neural networks parametrized by  $\gamma$ . This model is called a structured variational autoencoder (SVAE) [2].


**But what does it do?**

First of all, a neural network is neither conjugate nor exponential.  
How do we perform inference?

# Inference in SVAE

First of all, a neural network is neither conjugate nor exponential.  
How do we perform inference?


**Strategy:** Run VMP, but use “fake” messages for the neural network observation model.



# Inference in SVAE

First of all, a neural network is neither conjugate nor exponential.  
How do we perform inference?


**Strategy:** Run VMP, but use “fake” messages for the neural network observation model.



# Inference in SVAE

First of all, a neural network is neither conjugate nor exponential.  
How do we perform inference?

**Strategy:** Run VMP, but use “fake” messages for the neural network observation model.



# Messages in SVAE

The message from a **non-conjugate, non-exponential family** observation  $X$  to a parent  $Y$  is

$$m_{X \rightarrow Y} = r_{\xi}(t_X(X))$$

where  $r_{\xi}$  is a neural network whose output is the same shape of a message from a conjugate-exponential child.

# Messages in SVAE

The message from a **non-conjugate, non-exponential family** observation  $X$  to a parent  $Y$  is

$$m_{X \rightarrow Y} = r_\xi(t_X(X))$$


where  $r_\xi$  is a neural network whose output is the same shape of a message from a conjugate-exponential child.

## Inference in a SVAE

1. For a data  $\tau = \{s_0, a_0, s_1, a_1, \dots, s_T\}$ , perform VMP using SVAE messages
2. Compute the ELBO  $\mathcal{L}[q(\{x_i\}_{i=1}^T, \mu_\rho, \Sigma_\rho, \mathbf{F}, \Sigma)]$
3. Update neural networks with  $\nabla_{\gamma, \xi} \mathcal{L}[q(\{x_i\}_{i=1}^T, \mu_\rho, \Sigma_\rho, \mathbf{F}, \Sigma)]$
4. Update global parameters  $(\mu_\rho, \Sigma_\rho, \mathbf{F}, \Sigma)$  with natural gradients

# Autoencoding


How is a SVAE an *autoencoder*?





# Autoencoding


How is a SVAE an *autoencoder*?



We have two neural networks:

# Autoencoding

How is a SVAE an *autoencoder*?

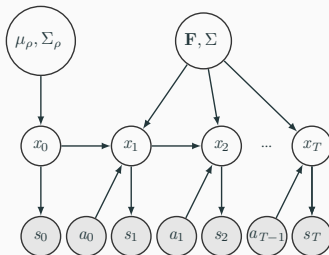


We have two neural networks:

1.  $r_\xi(s)$ : recognition network, takes real data and “encodes” it

# Autoencoding

How is a SVAE an *autoencoder*?



We have two neural networks:

1.  $r_\xi(s)$ : recognition network, takes real data and “encodes” it
2.  $\mu_\gamma(x), \Sigma_\gamma(x)$ : observation network, takes latent data and “decodes” it

# Conclusion

---

## Background

- Probabilistic graphical models

- Model learning

- Variational inference

## Structured variational autoencoder

- Model learning

## Conclusion

- Applications


- Current work

## Why SVAE?

The SVAE naturally applies to scenarios where there is already a tractable PGM.


# Why SVAE?

The SVAE naturally applies to scenarios where there is already a tractable PGM.



# Why SVAE?

The SVAE naturally applies to scenarios where there is already a tractable PGM.



In this scenario, the SVAE enables modeling non-Gaussian cluster shapes [2].

One idea I am currently working on is using the SVAE to learn latent models to be used in reinforcement learning<sup>1</sup>.

---

<sup>1</sup>Joint work with Marvin Zhang from UC Berkeley



One idea I am currently working on is using the SVAE to learn latent models to be used in reinforcement learning<sup>1</sup>.

## General approach:

- Learn a latent LDS
- Use MPC to control an agent in the latent space

---

<sup>1</sup>Joint work with Marvin Zhang from UC Berkeley

One idea I am currently working on is using the SVAE to learn latent models to be used in reinforcement learning<sup>1</sup>.

## General approach:

- Learn a latent LDS
- Use MPC to control an agent in the latent space

## Benefits of this approach:

---

<sup>1</sup>Joint work with Marvin Zhang from UC Berkeley

One idea I am currently working on is using the SVAE to learn latent models to be used in reinforcement learning<sup>1</sup>.

## General approach:

- Learn a latent LDS
- Use MPC to control an agent in the latent space

## Benefits of this approach:

- We can learn simple dynamics even with camera data

---

<sup>1</sup>Joint work with Marvin Zhang from UC Berkeley

One idea I am currently working on is using the SVAE to learn latent models to be used in reinforcement learning<sup>1</sup>.

## General approach:

- Learn a latent LDS
- Use MPC to control an agent in the latent space

## Benefits of this approach:

- We can learn simple dynamics even with camera data
- Model-based RL tends to be more sample efficient

---

<sup>1</sup>Joint work with Marvin Zhang from UC Berkeley

One idea I am currently working on is using the SVAE to learn latent models to be used in reinforcement learning<sup>1</sup>.

## General approach:

- Learn a latent LDS
- Use MPC to control an agent in the latent space

## Benefits of this approach:

- We can learn simple dynamics even with camera data
- Model-based RL tends to be more sample efficient

---

<sup>1</sup>Joint work with Marvin Zhang from UC Berkeley

One idea I am currently working on is using the SVAE to learn latent models to be used in reinforcement learning<sup>1</sup>.

## General approach:

- Learn a latent LDS
- Use MPC to control an agent in the latent space

## Benefits of this approach:

- We can learn simple dynamics even with camera data
- Model-based RL tends to be more sample efficient

## Demo

---

<sup>1</sup>Joint work with Marvin Zhang from UC Berkeley

Questions?

## References

---

- [1] J. Winn and C. Bishop. Variational message passing. *JMLR*, 2005.
- [2] M. Johnson, D. Duvenaud, A. Wiltchko, S. Datta, and R. Adams. Composing graphical models with neural networks for structured representations and fast inference. In *NIPS*, 2016.